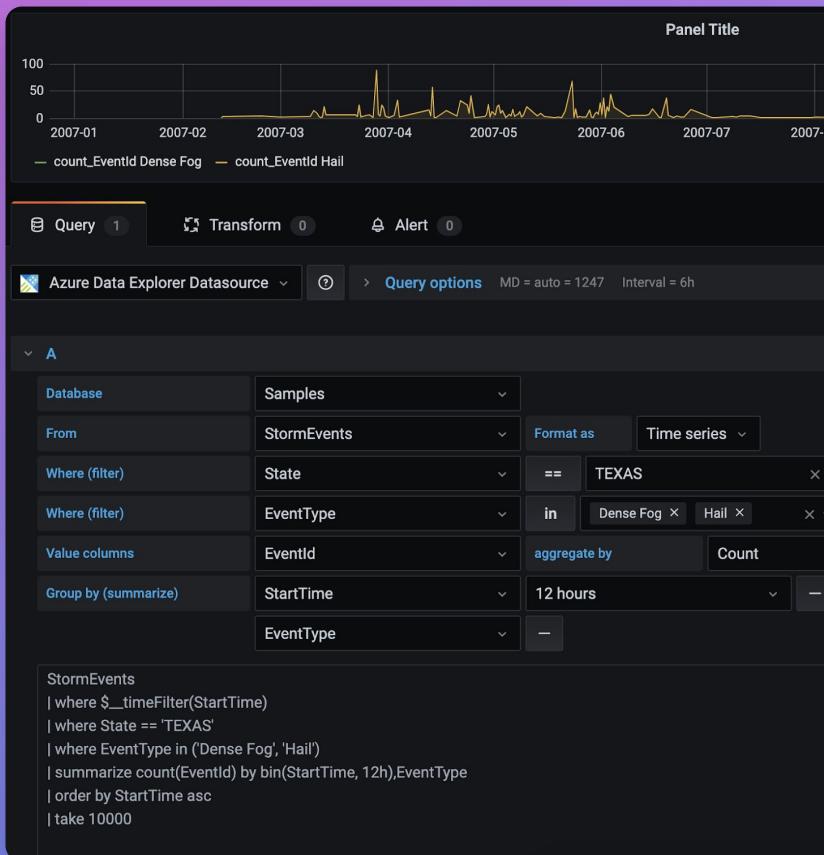




OpenTelemetry Collector deep dive



Juraci Paixão Kröhling
Software engineer
@jpkrohling



Presenter



Juraci Paixão Kröhling
Software engineer

Agenda

The basics

- What's OpenTelemetry
- What's OpenTelemetry Collector
- Other related projects: contrib and builder

Deployment patterns

- General purpose patterns (basic, normalizer, per-signal)
- Patterns for Kubernetes (daemonsets, sidecars)
- Enterprise patterns (multi-cluster, multitenant, load balancing)

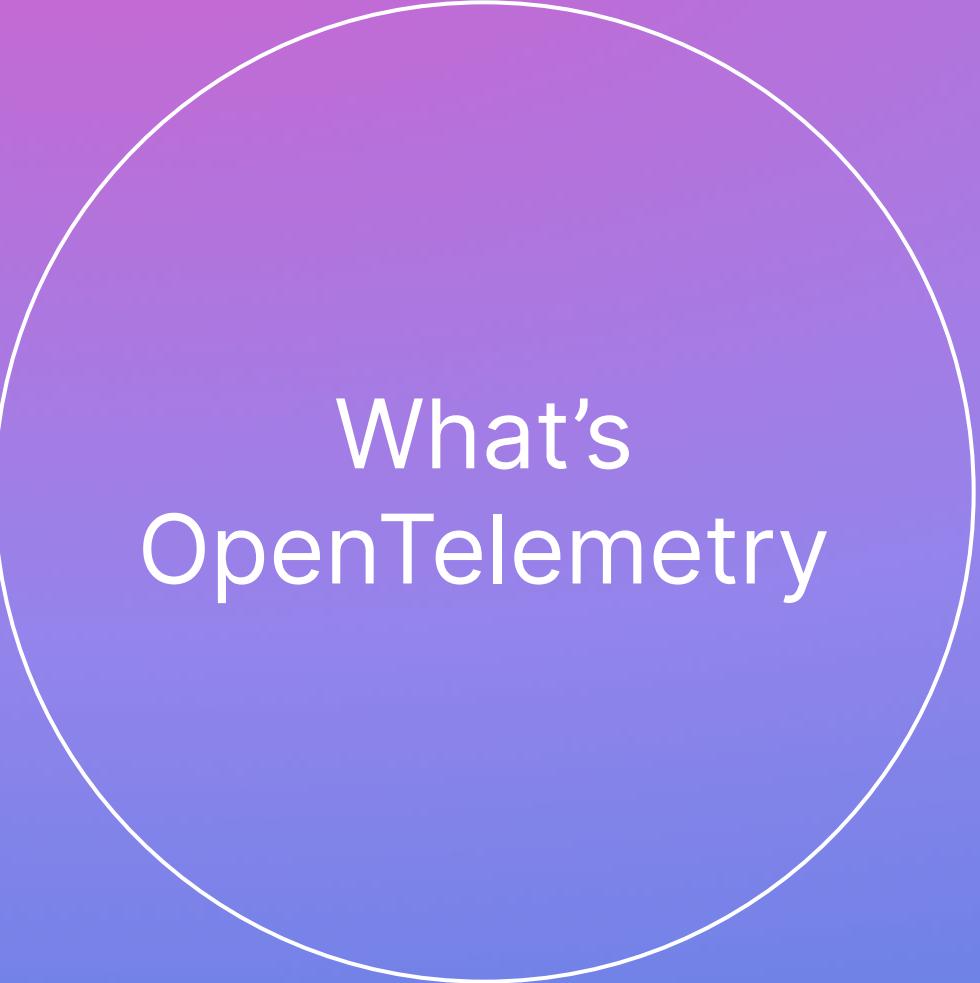
Advanced topics

- Assembling your own distribution
- Extending with your components

Questions and answers

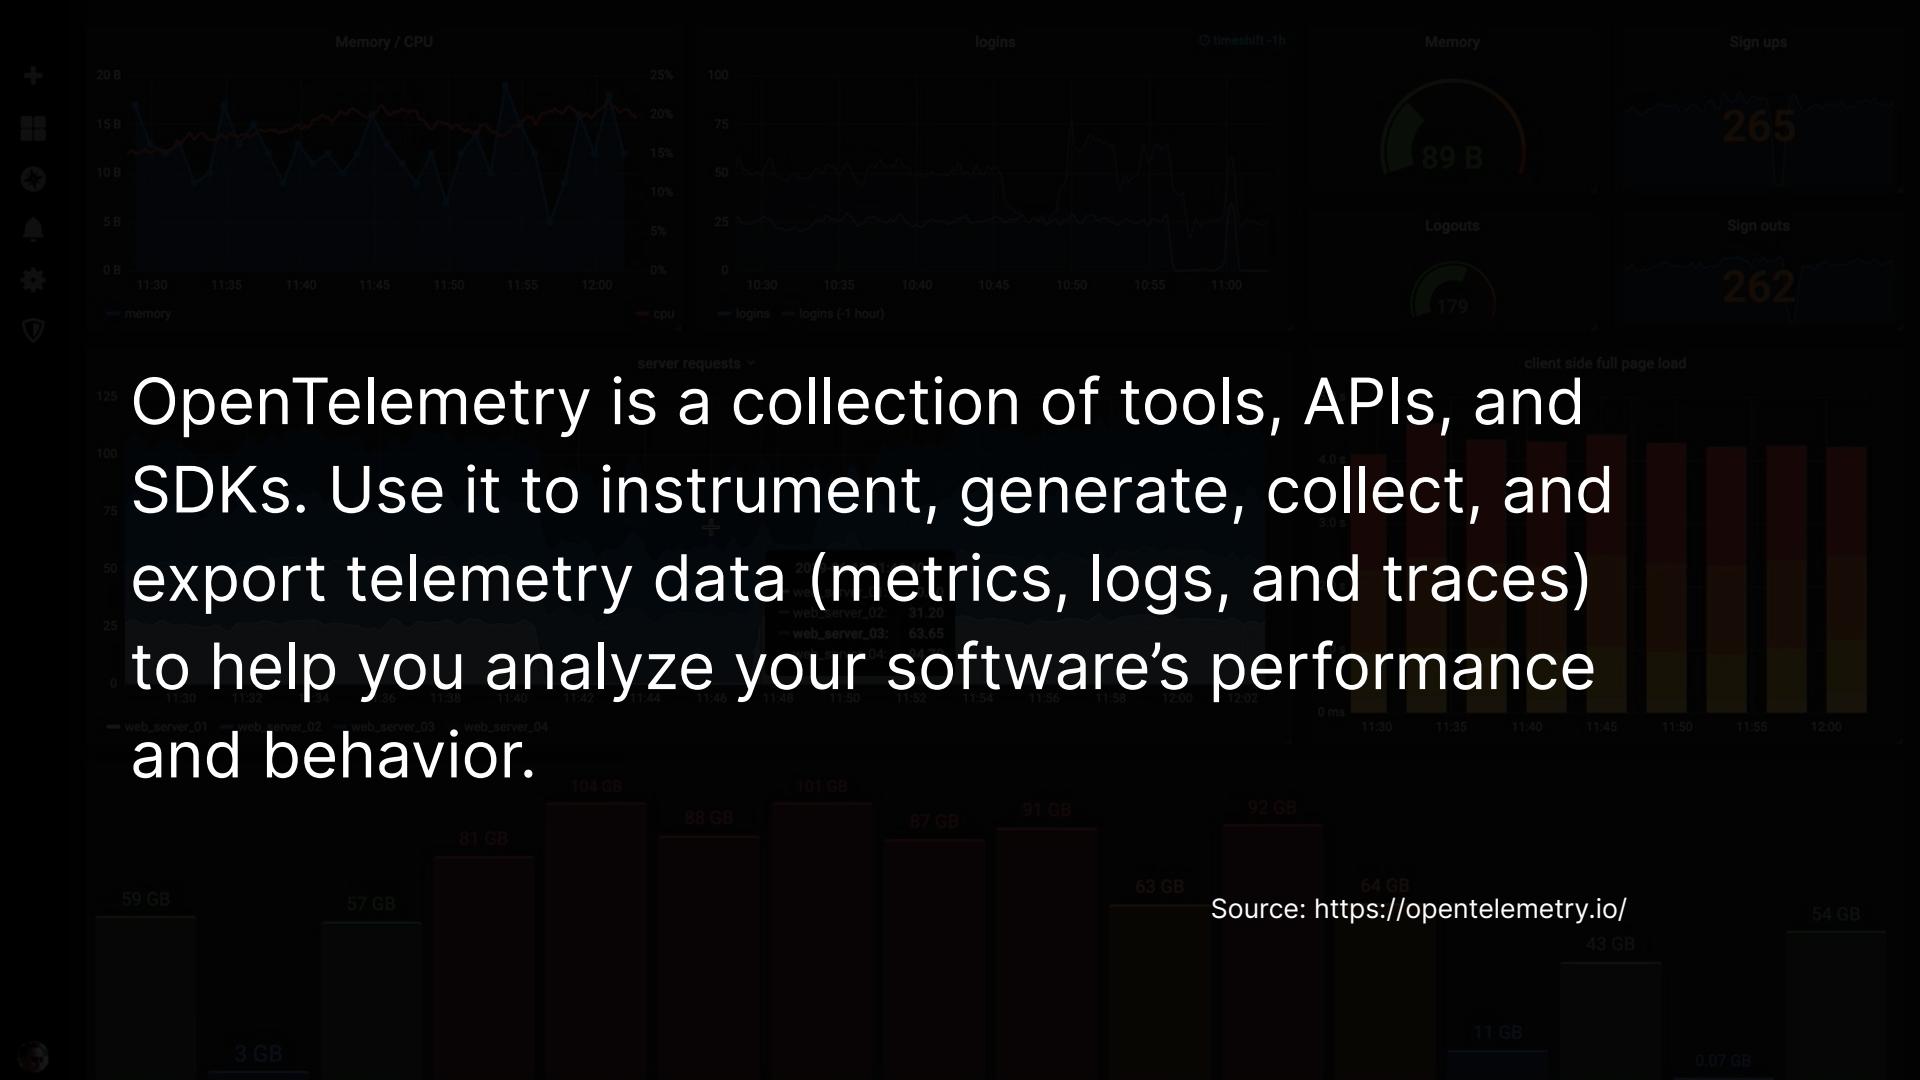


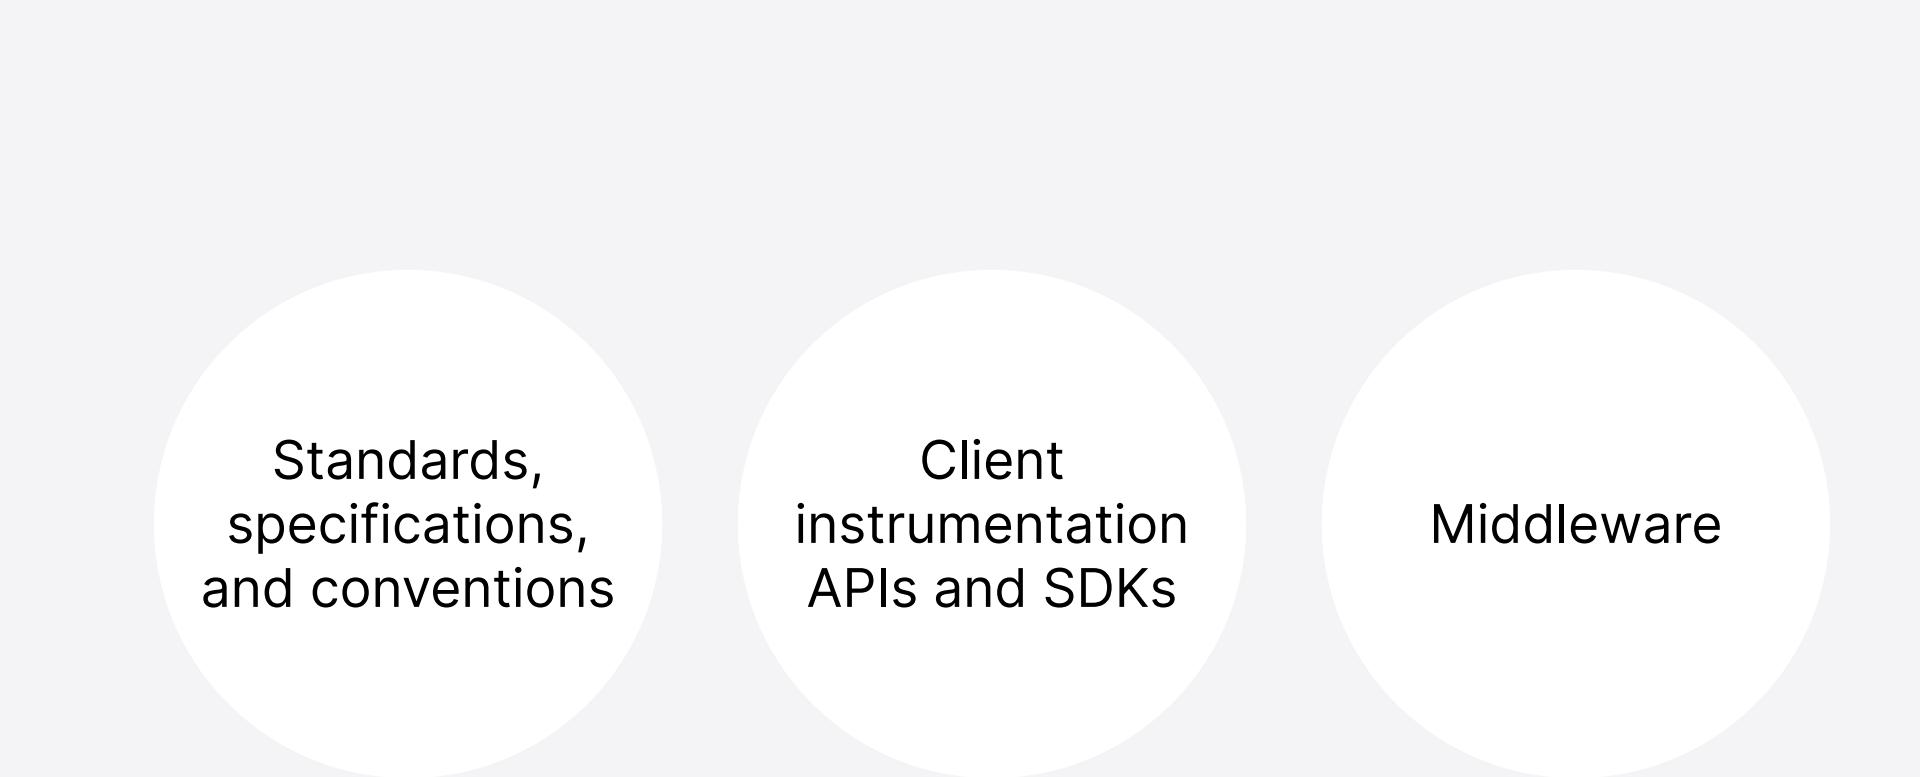
@jpkrohling



What's
OpenTelemetry







Standards,
specifications,
and conventions

Client
instrumentation
APIs and SDKs

Middleware



Standards, specifications, and conventions

OpenTelemetry API / SDK

- ✓ For people looking to implement APIs and SDKs

Semantic conventions

- ✓ Which metadata to include in which operations

OpenTelemetry Line Protocol

- ✓ Interface description language (IDL), specifying how data should look like and what endpoints should implement



Client instrumentation APIs and SDKs

OpenTelemetry API

- ✓ What you use on a daily basis to instrument your service

OpenTelemetry SDK

- ✓ What to do with the instrumentation: how to create the data, buffer, send out

Instrumentation libraries

- ✓ Libraries that will hook into parts of your stack and instrument it





Middleware



What's
OpenTelemetry
Collector



66

*Vendor-agnostic way to receive, process
and export telemetry data.*

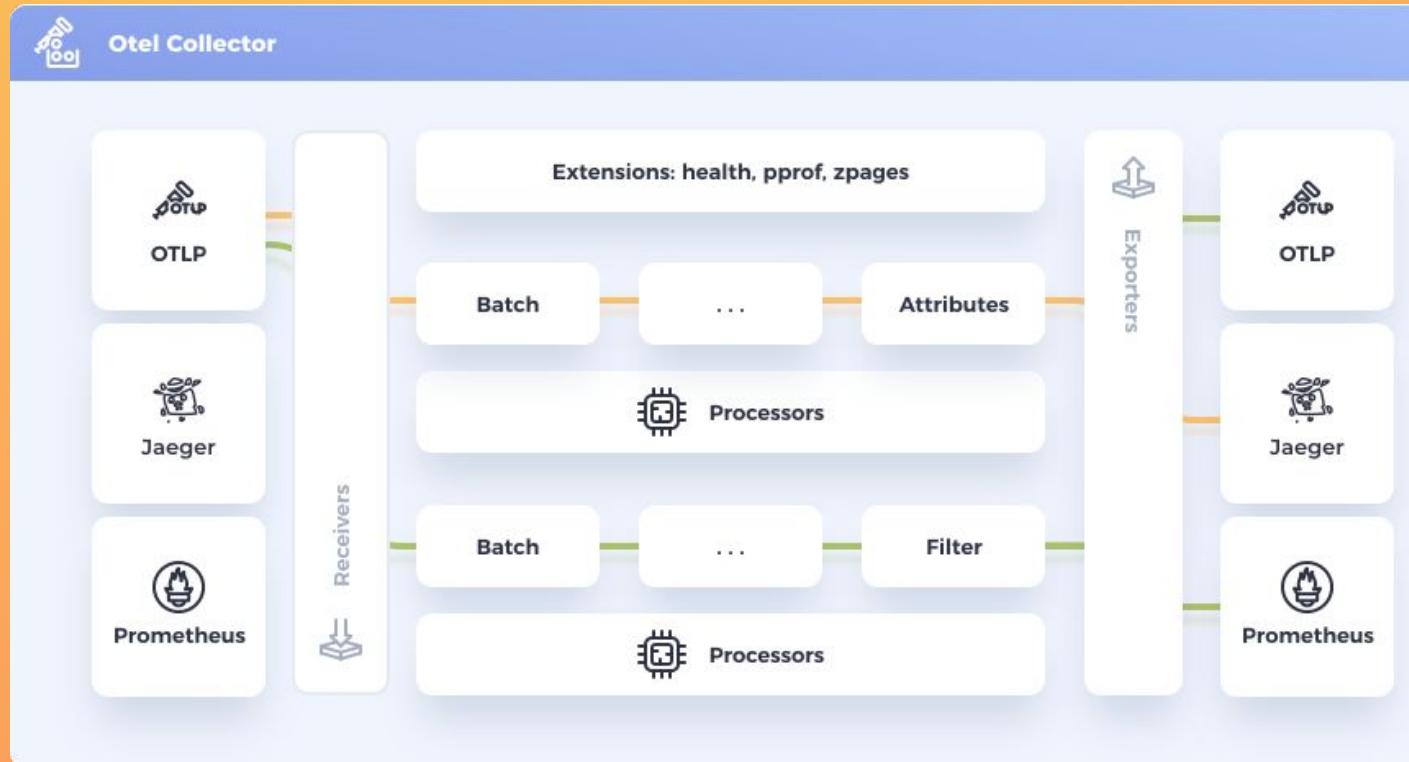


Source: <https://opentelemetry.io/docs/collector/>



@jpkrohling

OpenTelemetry Collector conceptual architecture



OpenTelemetry Collector - related projects

Contrib

- ✓ Where non-core components reside, such as vendor-specific ones

Builder

- ✓ Helper CLI tool to build OpenTelemetry Collector distributions

Operator

- ✓ Kubernetes operator managing OpenTelemetry Collector instances

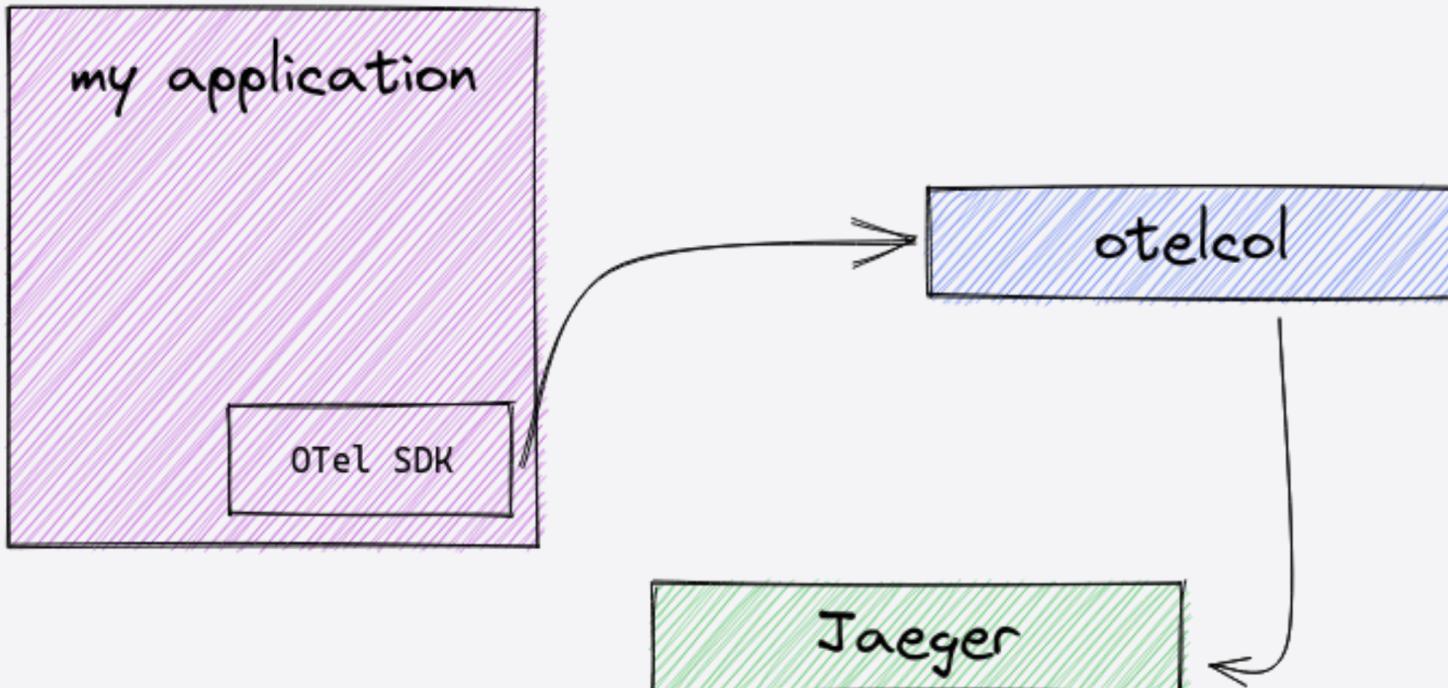




Patterns!



Pattern #1 - Basic I



Pattern #1 - Basic I

 Good for:

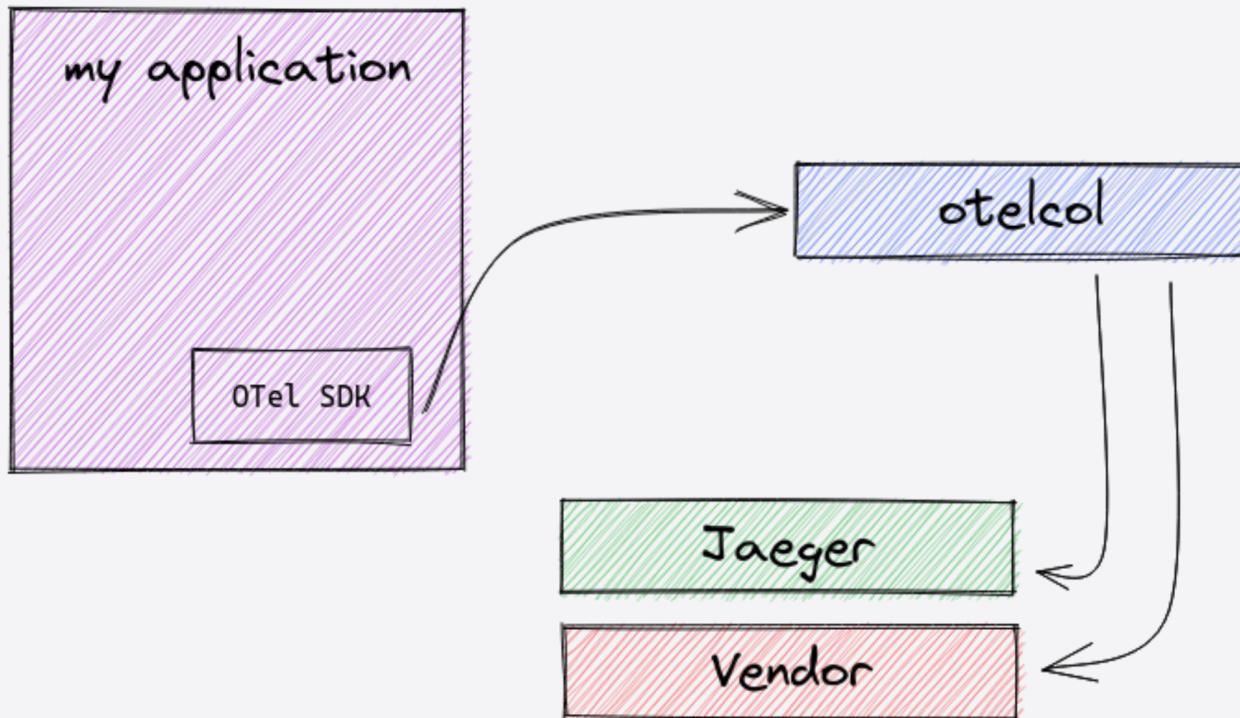
- Abstracting where to actually send the telemetry data
- Doing extra-processing between your workload and the telemetry backend

 Avoid when:

- Well, when you don't need an extra processing layer, every extra hop is a chance for things to go wrong 



Pattern #1 - Basic II - Fanout



Pattern #1 - Basic II - Fanout

Good for:

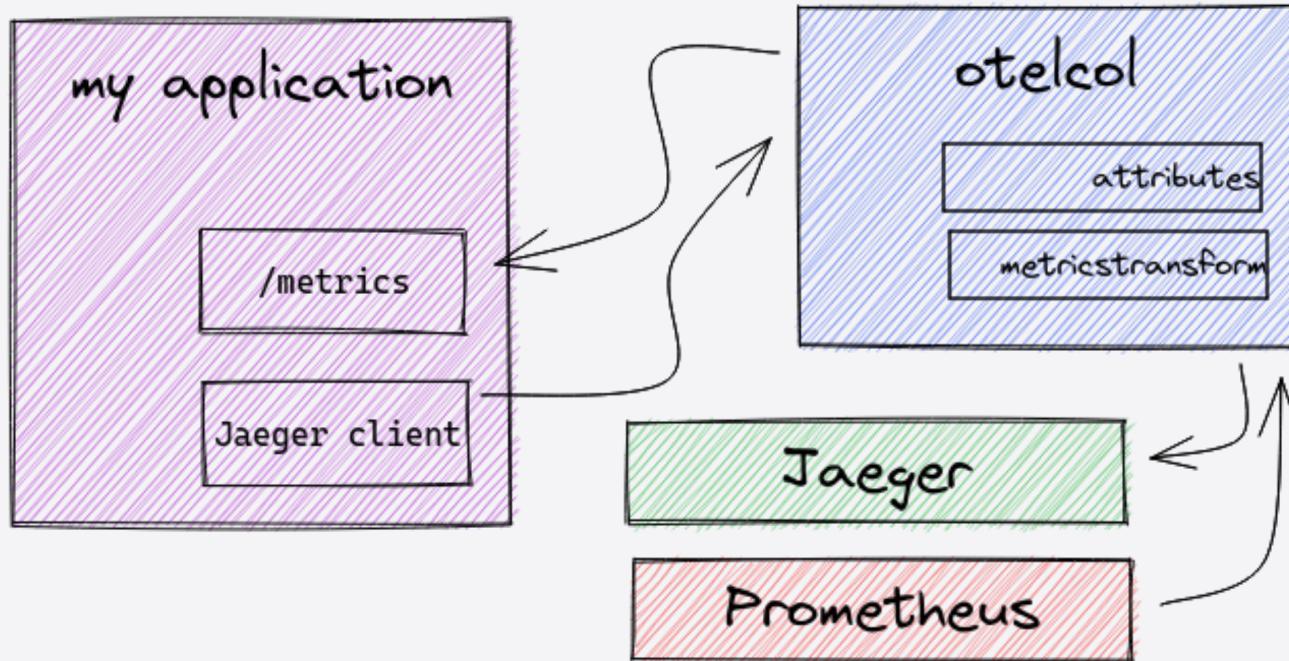
- Trying out different open source solutions and/or vendors
- Retaining data ownership even when your main observability tool is a SaaS

Avoid when:

- Processing tons of data: be conscious of the costs 



Pattern #2 - Normalizer



Pattern #2 - Normalizer

Good for:

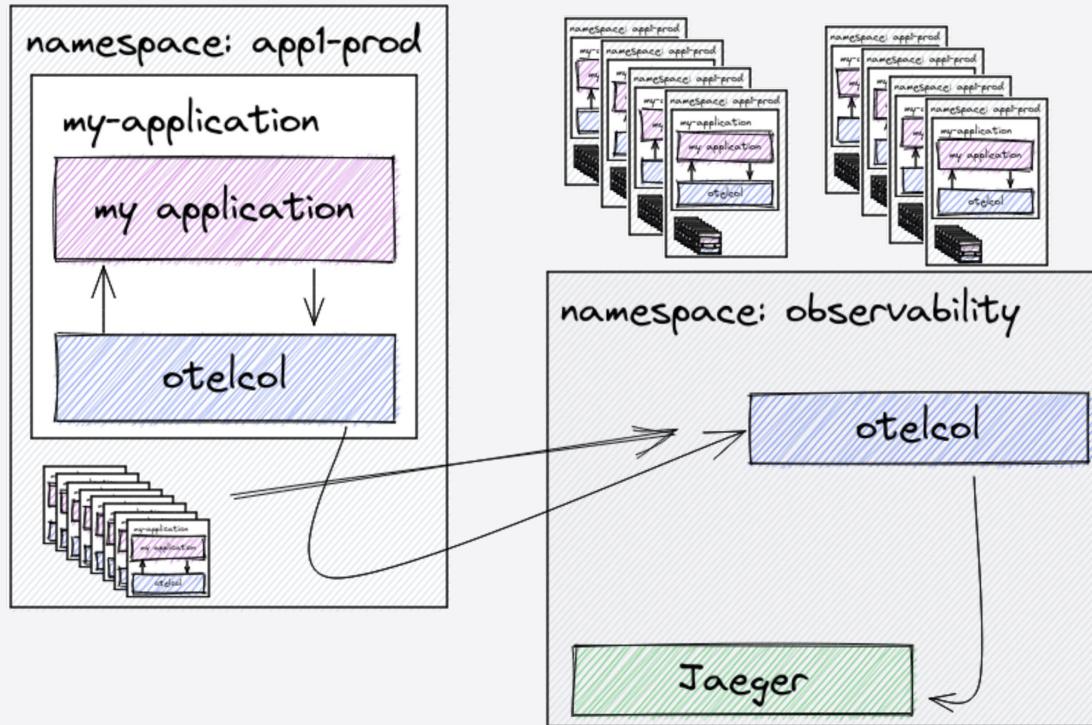
- Ensuring that different data points have the same semantics for the same things
- It's hard or undesirable to fix the problem at the source

Avoid when:

- You have too many things to normalize. It might be better to try to  fix the problem at the source



Pattern #3 - Kubernetes - Sidecars



Pattern #3 - Kubernetes - Sidecars

Good for:

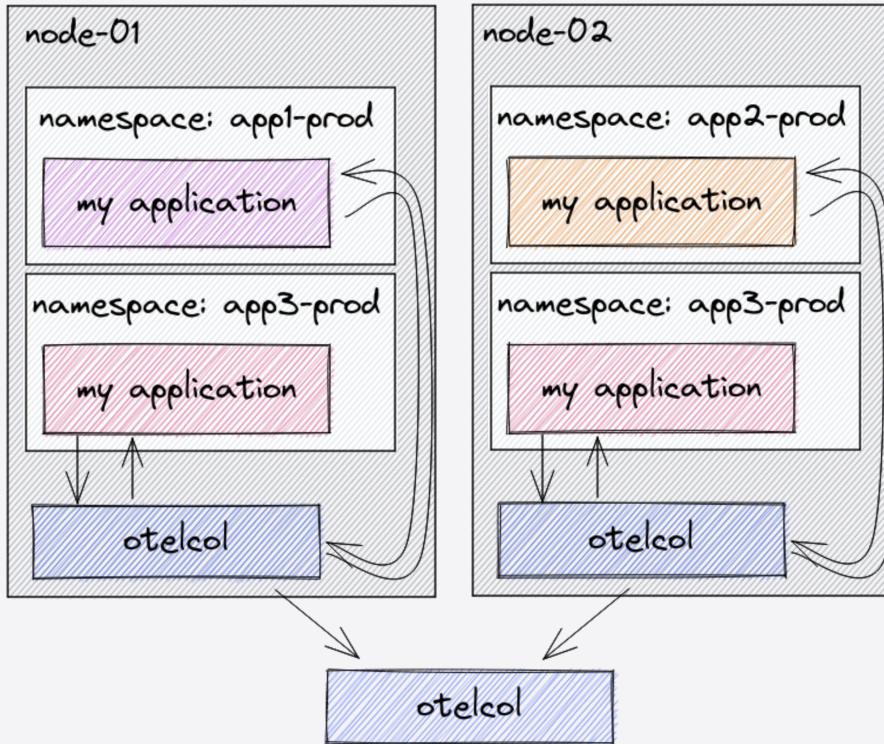
- Quickly offloading telemetry data from your application to a local process
- Fine-grained control over the configuration for each PodSpec or namespace
- Client-side load balancing is better when there are multiple of clients, especially for long-lived connections (HTTP/2, gRPC, Thrift, ...)

Avoid when:

- The overhead is not acceptable, as each sidecar needs at least ~20MiB of RAM
- You can't use something like the operator to manage the configs



Pattern #3 - Kubernetes - DaemonSets



Pattern #3 - Kubernetes - DaemonSets

Good for:

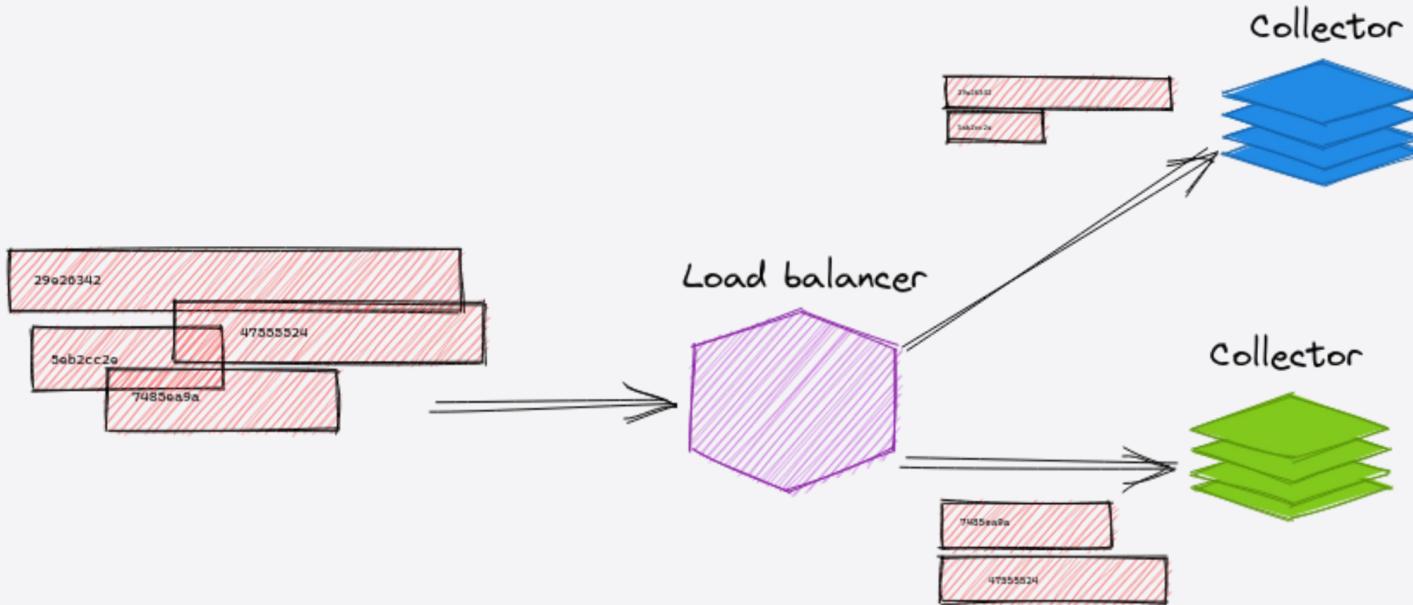
- Quickly offloading telemetry data from your application to a local process
- Less collector instances mean less maintenance and runtime overhead

Avoid when:

- You need multi-tenancy
- It's not acceptable to lose telemetry data for all pods on a node in case of a  crash with the local collector



Pattern #4 - Load balancing



Pattern #4 - Load balancing

Good for:

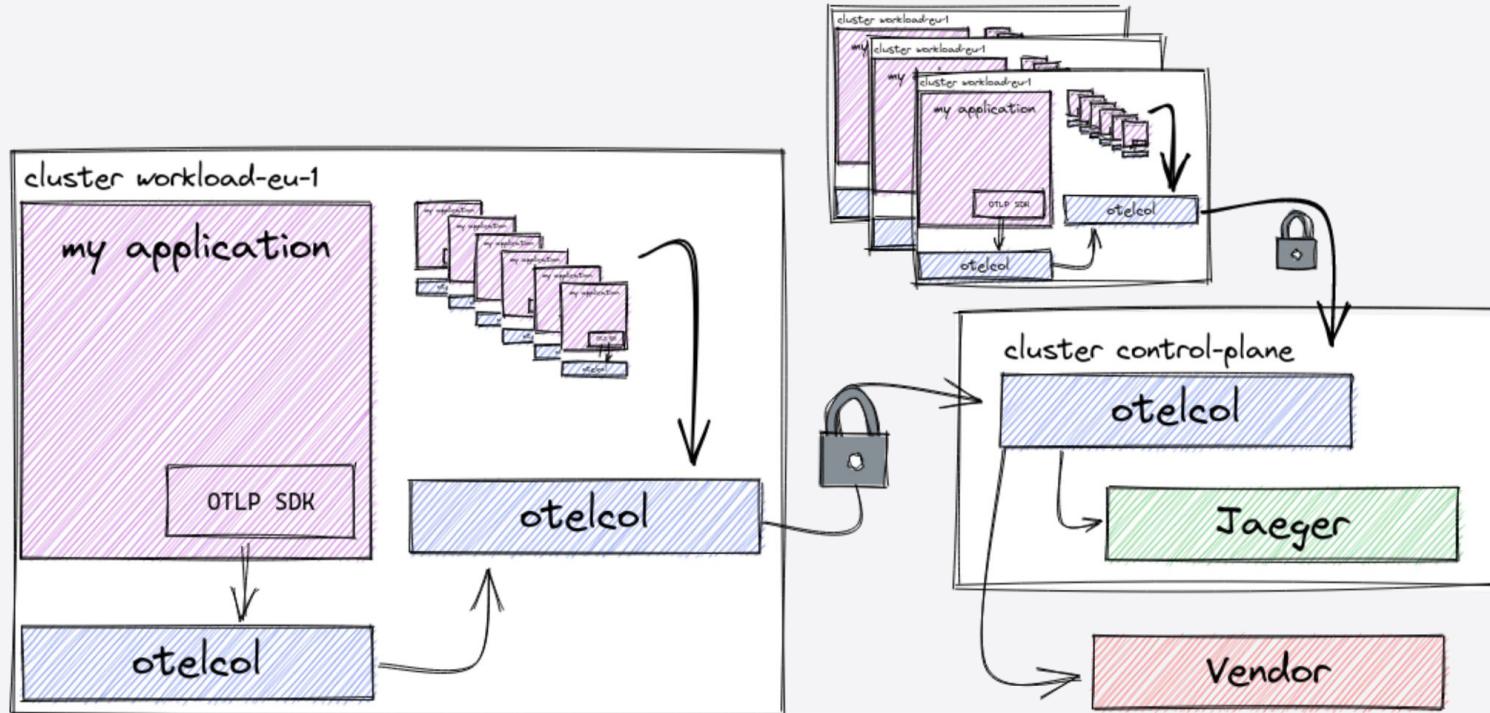
- Load balancing whole traces to collectors that need a complete view of the trace:
span metrics processor, tail-based sampling, ...

Avoid when:

- You just need a simple load balancing, without caring about the trace ID at all. For that, use a regular HTTP/2 or gRPC load balancer.



Pattern #5 - Multi-cluster



Pattern #5 - Multi-cluster

Good for:

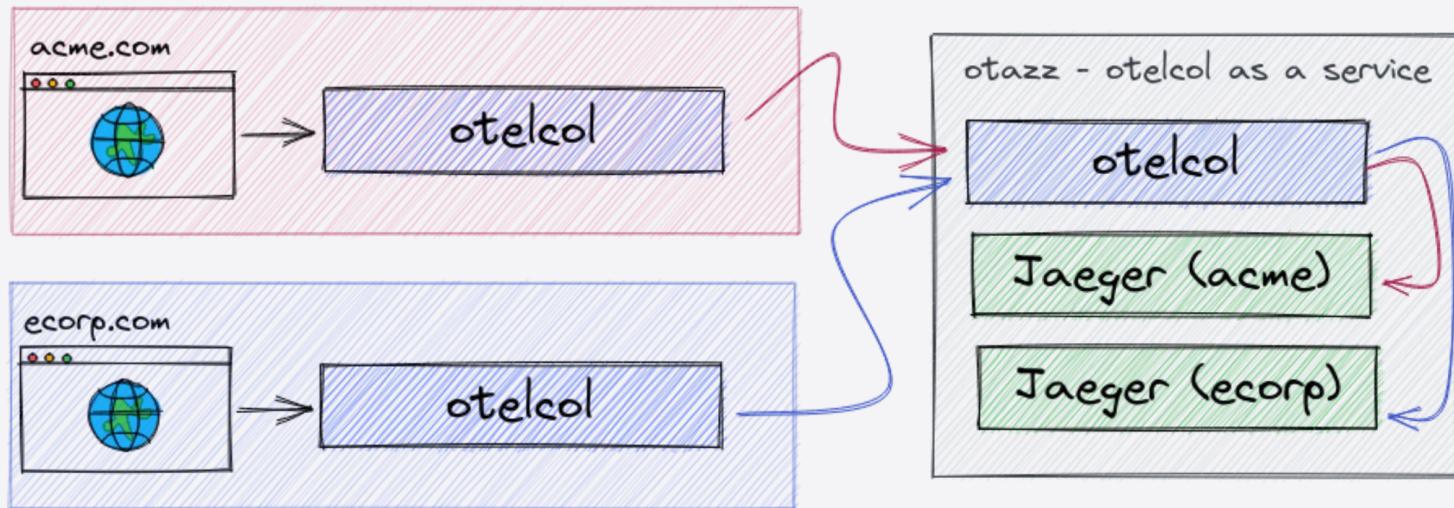
- Centralizing your telemetry data collection across clusters
- Running business analytics on all of your telemetry data

Avoid when:

- You can have your control plane to query data directly on individual clusters
- Networking costs are a concern



Pattern #6 - Multitenant



Pattern #6 - Multitenant

Good for:

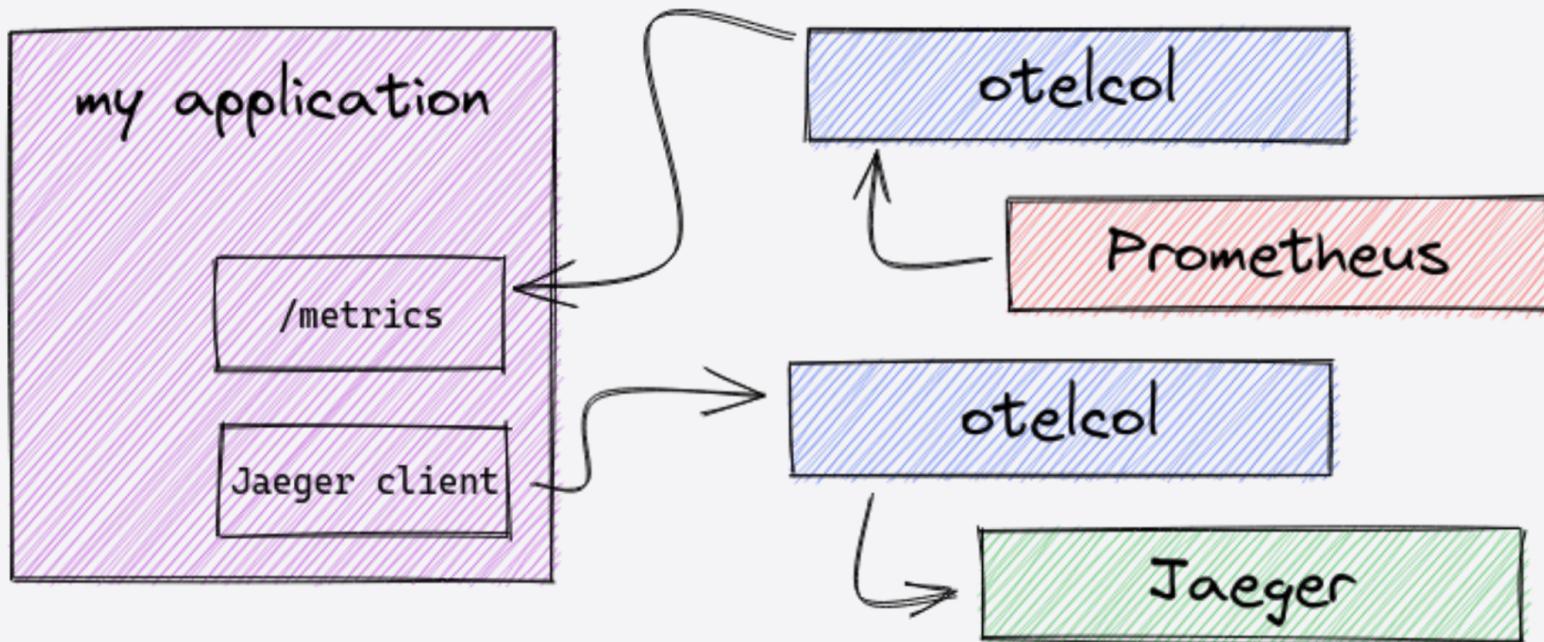
- Small deployments, where a central collector processes all the telemetry data for all tenants
- Central teams to handle telemetry backends for multiple departments

Avoid when:

- You can have one entrypoint per tenant, avoiding a single point of failure



Pattern #7 - Per signal



Pattern #7 - Per signal

 Good for:

- Isolating failures on production environments

 Avoid when:

- You just need a simple deployment for your local dev or staging environments



Customizing OpenTelemetry Collector



Building your own distribution

1. manifest.yaml
2. opentelemetry-collector-builder
3. ???
4. Profit!



Building a component

- Config
- The component code, implementing one or more interfaces
- Factory



Building a processor

- Bootstrap go module
- Create a config.go
- Create a processor.go
- Add the processor logic
- Create a factory.go
- Bonus points: metrics.go





Key takeaways



Key takeaways

- OpenTelemetry has different subprojects in different areas
- The collector works as a middleware, abstracting the telemetry backends from your workloads
- It has tons of components for you to experiment with
- Mix and match collector instances with potentially different configurations
- Use the patterns from this presentation to derive your own patterns
- Extending the collector isn't that hard!
- Building your own distribution might be a good idea depending on your use cases



Thanks for attending!

Have more questions?

@jpkrohling at twitter and github

Get involved:



#otel-collector
(CNCF Slack)



open-telemetry/
opentelemetry-collector